

# Operating Systems (Client)

## Lecture 4 Uniprocessor Scheduling

Dr. Kevin Farrell

# **Uniprocessor Scheduling**

# Goals of Scheduling

- Quick response time
- Fast throughput
- Processor efficiency
- To be fair to all users
- Degrade performance gracefully
- Be consistent and predictable

# Type of Scheduling

## ■ Long-term

- performed when new process is created

## ■ Medium-term → Medium Level

- swapping (Virtual Memory)

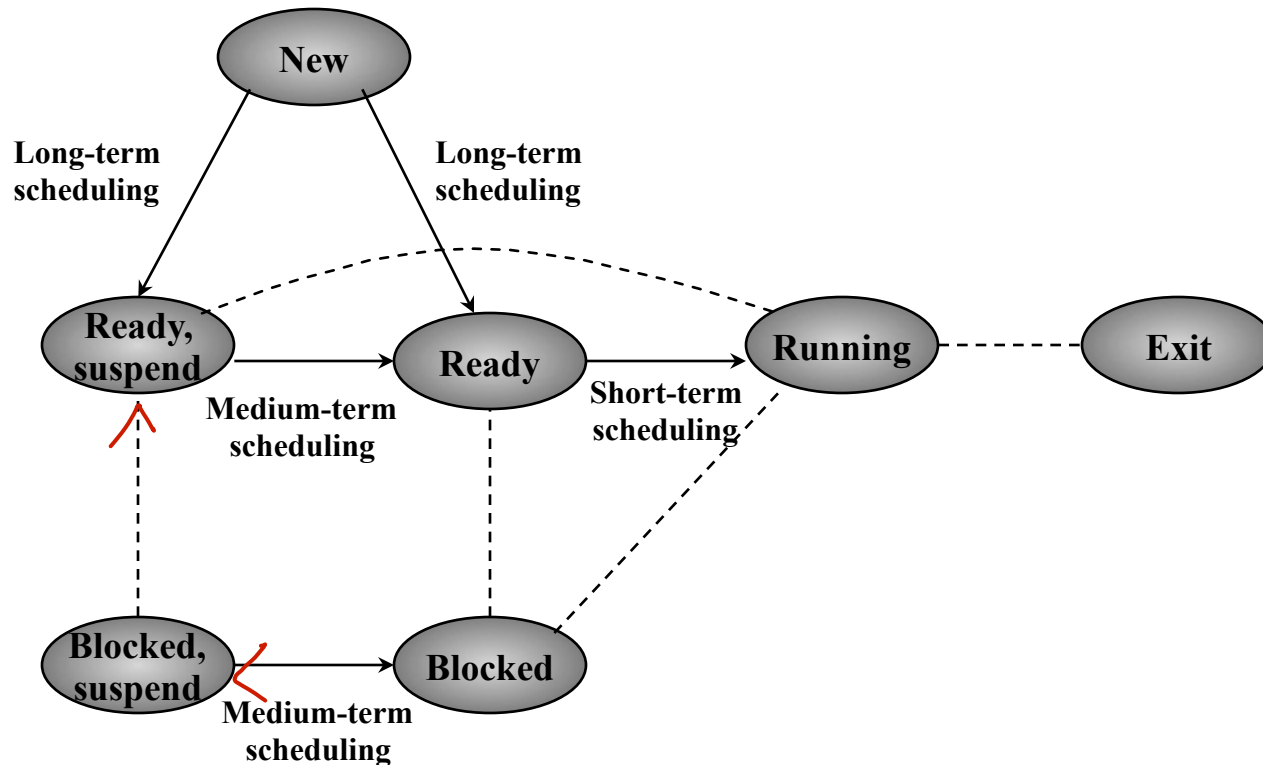
## ■ Short-term → Low Level

- which ready process to execute next

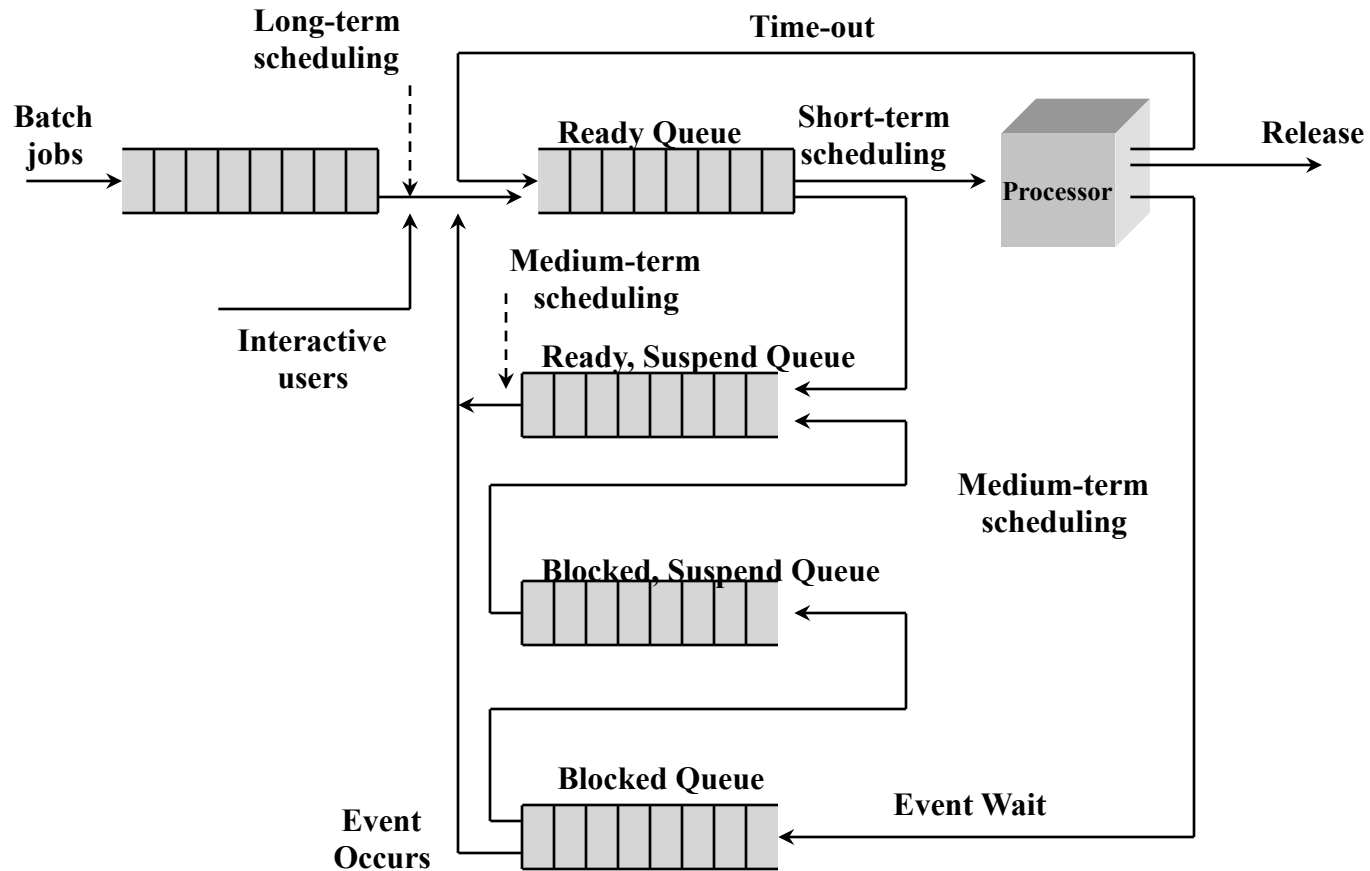
## ■ I/O

- decision as to which process's pending I/O request shall be handled by available I/O device

# Scheduling and Process State Transition



# Queuing Diagram for Scheduling




# Long-Term Scheduling

- Determines which programs are admitted to the system for processing
- Controls the degree of multiprogramming
- More processes, smaller percentage of time each process is executed

# Medium-Term Scheduling

- Swapping
- Based on the need to manage multiprogramming



transfer of pages  
between physical and  
virtual memory



# Short-Term Scheduling

- Sometimes known as the dispatcher (but the dispatcher is just a portion of the short-term scheduler)
- Invoked when an event occurs
  - clock interrupts
  - I/O interrupts
  - operating system calls
  - signals

# Short-Term Scheduling Criteria

## ■ User-oriented

### ■ Response Time

- Elapsed time between the submission of a request until there is output.

## ■ System-oriented

- effective and efficient utilization of the processor

# Short-Term Scheduling Criteria

- Performance-related
  - measurable such as response time and throughput
- Not performance related
  - predictability

Need to use a  
'metric' to measure  
performance.

# Priorities

Execute next, or, put to  
the head of the Queue

- Scheduler will always choose a process of higher priority over one of lower priority
- Have multiple ready queues to represent each level of priority — 'Multi-level Queues'
- Lower-priority may suffer starvation
  - allow a process to change its priority based on its age or execution history

Multi-level  
feedback Queues

may never be given an opportunity  
to be executed in a reasonable time

# Decision Mode

*Generally, used in older systems*

## ■ Nonpreemptive

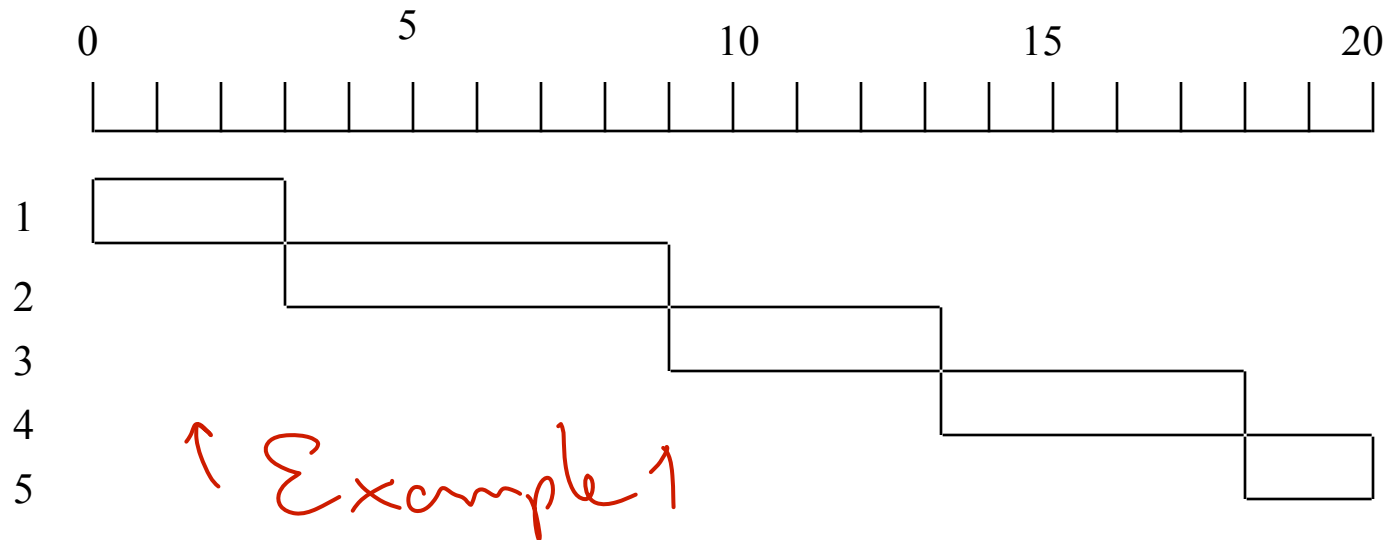
- Once a process is in the running state, it will continue until it terminates or blocks itself for I/O

## ■ Preemptive

*in modern systems*

- Currently running process may be interrupted and moved to the Ready state by the operating system
- Allows for better service since any one process cannot monopolize the processor for very long

# First-Come-First-Served (FCFS)



- Each process joins the Ready queue
- When the current process ceases to execute, the oldest process in the Ready queue is selected

# First-Come-First-Served (FCFS)

- A short process may have to wait a very long time before it can execute
- Favors CPU-bound processes
  - I/O-bound processes have to wait until CPU-bound process complete

## Example 2

'metric 1'


### Example of FCFS Policy: (p67 Ritchie)

	(a)	(b)	(c)
<u>Job</u>	<u>Est. Run Time</u>	<u>Waiting</u>	<u>Ratio b/a</u>
1	2	0	0
2	60	2	$\frac{2}{60} = 0.03$
3	1	62	62
4	3	63	21
5	50	66	1.32

'Metric 2' = Average Wait time =  $193 \div 4 = 48.25$  (approx.)  
(Exclude the first process)



# Shortest Job First (SJF)



- Nonpreemptive policy
- Process with shortest expected processing time is selected next
- Short process jumps ahead of longer processes

marks overhead

# Shortest Job First


- Predictability of longer processes is reduced
- If estimated time for process not correct, the operating system may abort it
- Possibility of starvation for longer processes

## Example 2

	(a)	(b)	(c)
<i>Job</i>	<i>Est. Run Time</i>	<i>Waiting</i>	<i>Ratio b/a</i>
3	1	0	0.0
1	2	1	0.5
4	3	3	1.0
5	50	6	0.1
2	60	56	0.9

Average Wait time =  $66 \div 4 = 16.5$  (approx)  
Recall FCFS Av. wait time = 48s !!!


# Shortest Remaining Time



- Preemptive version of shortest process next policy

- Must estimate processing time  $\Rightarrow$  incurs overhead


# Highest Response Ratio Next (HRRN)



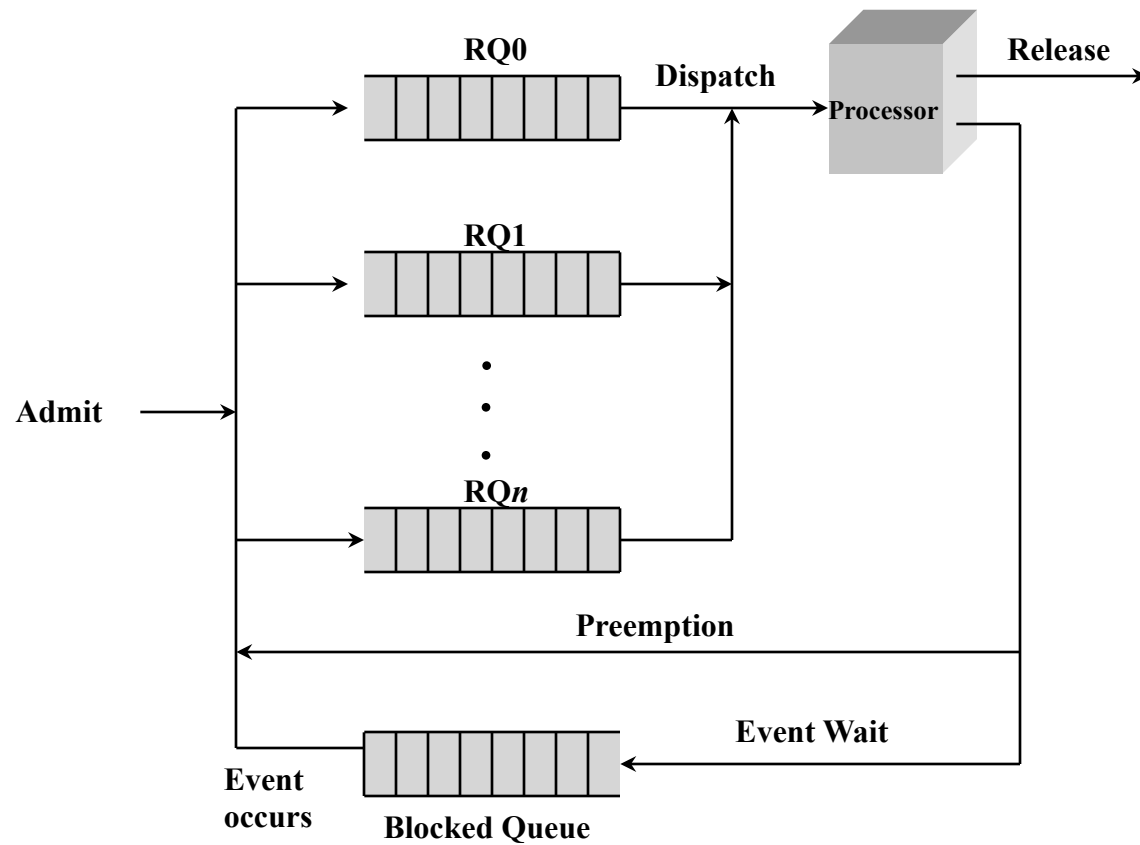
■ Choose next process with the highest ratio

$$P = \frac{\text{time waiting} + \text{expected run time}}{\text{expected run time}}$$


$$P = t_{\text{wait}} \left( \frac{1}{t_{\text{exp}}} \right) + 1$$




# Priority Queuing



# Round-Robin



- Uses preemption based on a clock
- An amount of time is determined that allows each process to use the processor for that length of time: **Time Quantum**





(a) Job A

0

8

time quantum = 10

(b) Job A

Job A

0

5

8

time quantum = 5

(c)

Job A

Job A

Job A

Job A

Job A

Job A

Job A

Job A

0

1

2

3

4

5

6


7

8


time quantum = 1

# Multi-Level Feedback Queues

*Pre-emptable Queues*



# Feedback



- Penalize jobs that have been running longer
- Don't know remaining time process needs to execute

# Fair Scheduling

- User's application runs as a collection of processes (threads)
- User is concerned about the performance of the application
- Need to make scheduling decisions based on groups of processes

# UNIX Scheduling

- Priorities are recomputed once per second
- Base priority divides all processes into fixed bands of priority levels
- “Nice” adjustment factor used to keep process in its assigned band (cf. “To renice a process”)

# Feedback

- Process is demoted to the next lower-priority queue each time it returns to the ready queue
- Longer processes drift downward
- To avoid starvation, preemption time for lower-priority processes is longer

# References

- “Operating Systems - Internals and Design Principles” - William Stallings (Prentice Hall, 4th edition, 2000)