

Operating Systems(Server)

Lecture 5 Process Synchronisation 2

Dr. Kevin Farrell

Chapter 7 Part 2: Process Synchronization 2

- Semaphores
- Classical Problems of Synchronization
- Critical Regions
- Monitors
- Synchronization in Solaris 2 & Windows 2000

Semaphores

- Synchronization tool for generalised solutions to more complex problems
- A Semaphore, S , is an integer variable
- Semaphores are low-level synchronisation mechanisms; i.e. the operations on semaphores are implemented as **system calls**
- **The operations below are the traditional (Dijkstra) definitions of wait(S) and signal(S)**
- Can only be accessed via two indivisible (atomic) operations:

```
wait ( $S$ )  
{  
    while( $S \leq 0$ )  
        ; // Do nothing  
     $S - -$ ;  
}
```

```
signal ( $S$ )  
{  
     $S++$ ;  
}
```

Semaphore Usage: Critical Section of n Processes

- Shared data:

semaphore mutex; //initially $mutex = 1$

- Process P_i with mutual exclusion as follows:

```
do {  
    wait(mutex);  
    critical section  
    signal(mutex);  
    remainder section  
} while (1);
```

Problems with this Semaphore Definition

- Main problem with previous definition of mutual exclusion solutions to critical-section problem, and with our current definition of *wait(S)* on the semaphore is that they all require **busy waiting**.
- Busy waiting wastes CPU cycles that some other process might be able to use productively.
- This type of semaphore is also called a **spinlock** because the process “spins” (loops) while waiting for the lock to be released.
- Spinlocks are useful in multiprocessor systems: suppose process, P_0 has the lock and is executing on one processor, then process, P_1 can spin on another processor, waiting for the lock, while P_0 continues to execute.
- Spinlocks are advantageous in situations where context switching to another READY process, P_2 (say), would take considerable time, whereas the spinlock of P_1 would only occupy the processor for a relatively short time
- i.e. if locks are expected to be held for short times, allow spinlocks, and therefore avoid having to make time-consuming context-switches

Semaphore: New Implementation

- **New Definition:** define a semaphore as a record

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```
- Assume two simple operations:
 - ✦ **block** suspends the process that invokes it.
 - ✦ **wakeup(*P*)** resumes the execution of a blocked process **P**.

New Implementation: Code

- Semaphore operations now defined as

wait(S):

S.value- -;

if (S.value < 0) {

 add this process to **S.L**;

block();

}

signal(S):

S.value++;

if (S.value <= 0) {

 remove a process **P** from **S.L**;

wakeup(P);

}

New Implementation: Atomicity

- We need to implement wait() and signal() as atomic operations
- This is a critical section problem in its own right!
- On Uniprocessor systems, this can be achieved by:
 - ◆ Disabling Interrupts
- On Multiprocessor systems, disabling interrupts is undesirable. Instead, use
 - ◆ Hardware synchronisation mechanisms (for eg.: TSL) if they exist, or
 - ◆ Software synchronisation mechanisms such as our critical section algorithms of Lecture 5.
- The new implementation of the Semaphore has the following two important consequences:
 - ◆ Once a process blocks itself, it sleeps in a wait-queue in a BLOCKED state => it doesn't use CPU cycles for waiting. Therefore, it eliminates busy-waiting from application program code.
 - ◆ There *is* some busy-waiting for short periods in the algorithms used to implement the wait() and signal() operations atomically.

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

- Let S and Q be two semaphores initialized to 1

P_0	P_1
$wait(S);$	$wait(Q);$
$wait(Q);$	$wait(S);$
\vdots	\vdots
$signal(S);$	$signal(Q);$
$signal(Q)$	$signal(S);$

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended. This would happen, for example, if we implemented a LIFO policy to remove processes from the waiting queue. (LIFO = Last In First Out)

Two Types of Semaphores

- *Counting* semaphore – integer value can range over an unrestricted domain.
- *Binary* semaphore – integer value can range only between 0 and 1; can be simpler to implement.
- Can implement a counting semaphore S as a binary semaphore.

Implementing *S* as a Binary Semaphore

- Data structures:
 - **binary-semaphore S1, S2;**
 - int C;**
- Initialization:
 - S1 = 1**
 - S2 = 0**
 - C = initial value of semaphore S**

Implementing *S*

- *wait* operation
 wait(S1);
 C--;
 if (C < 0) {
 signal(S1);
 wait(S2);
 }
 signal(S1);

- *signal* operation
 wait(S1);
 C ++;
 if (C <= 0)
 signal(S2);
 else
 signal(S1);

Classical Problems of Synchronization

- Bounded-Buffer Producer Consumer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded-Buffer Problem (i.e. Producer-Consumer Problem)

- Shared data:
 - ✦ semaphore full, empty, mutex;
- empty = semaphore giving #empty buffer positions
- full = semaphore giving #full buffer positions
- mutex = semaphore for providing mutual exclusion

Initially:

full = 0, empty = n, mutex = 1

Bounded-Buffer Problem Producer Process

```
do {  
    ...  
    produce an item in nextp  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```

Bounded-Buffer Problem Consumer Process

```
do {  
    wait(full)  
    wait(mutex);  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```


Readers-Writers Problem

- Shared data

semaphore mutex, wrt;

Initially

mutex = 1, wrt = 1, readcount = 0

Readers-Writers Problem Writer Process

wait(wrt);

...

writing is performed

...

signal(wrt);


Readers-Writers Problem Reader Process

```
wait(mutex);
readcount++;
if (readcount == 1)
    wait(wrt);
signal(mutex);

...
reading is performed

...
wait(mutex);
readcount--;
if (readcount == 0)
    signal(wrt);
signal(mutex);
```

Dining-Philosophers Problem



- Shared data
semaphore chopstick[5];
Initially all values are 1

Dining-Philosophers Problem

```
■ Philosopher i:  
  do {  
    wait(chopstick[i])  
    wait(chopstick[(i+1) % 5])  
    ...  
    eat  
    ...  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    ...  
    think  
    ...  
  } while (1);
```

Critical Regions

- High-level synchronization construct
- A shared variable ***v*** of type ***T***, is declared as:
v*: shared *T
- Variable ***v*** accessed only inside statement
region *v* when *B* do *S*

where ***B*** is a boolean expression.

- While statement ***S*** is being executed, no other process can access variable ***v***.

Critical Regions

- Regions referring to the same shared variable exclude each other in time.
- When a process tries to execute the region statement, the Boolean expression B is evaluated. If B is true, statement S is executed. If it is false, the process is delayed until B becomes true and no other process is in the region associated with v .

Example – Bounded Buffer

■ Shared data:

```
struct buffer {  
    int pool[n];  
    int count, in, out;  
}
```


Bounded Buffer Producer Process

- Producer process inserts **nextp** into the shared buffer

```
region buffer when( count < n) {  
    pool[in] = nextp;  
    in:= (in+1) % n;  
    count++;  
}
```

Bounded Buffer Consumer Process

- Consumer process removes an item from the shared buffer and puts it in **nextc**

```
region buffer when (count > 0) {                                nextc =  
pool[out];  
    out = (out+1) % n;  
    count--;  
}
```

Implementation of region x when B do S

- Associate with the shared variable x , the following variables:

semaphore mutex, first-delay, second-delay;
int first-count, second-count;

- Mutually exclusive access to the critical section is provided by **mutex**.
- If a process cannot enter the critical section because the Boolean expression **B** is false, it initially waits on the **first-delay** semaphore; moved to the **second-delay** semaphore before it is allowed to reevaluate B .

Implementation

- Keep track of the number of processes waiting on **first-delay** and **second-delay**, with **first-count** and **second-count** respectively.
- The algorithm assumes a FIFO ordering in the queuing of processes for a semaphore.
- For an arbitrary queuing discipline, a more complicated implementation is required.

Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.


- **monitor** *monitor-name*

```
{  
    shared variable declarations  
    procedure body P1 (...) {  
        . . .  
    }  
    procedure body P2 (...) {  
        . . .  
    }  
    procedure body Pn (...) {  
        . . .  
    }  
    {  
        initialization code  
    }  
}
```


Monitors

- To allow a process to wait within the monitor, a **condition** variable must be declared, as
condition x, y;
- Condition variable can only be used with the operations **wait** and **signal**.
 - ✦ The operation **x.wait();**
means that the process invoking this operation is suspended until another process invokes **x.signal();**
 - ✦ The **x.signal** operation resumes exactly one suspended process. If no process is suspended, then the **signal** operation has no effect.

Schematic View of a Monitor



Monitor With Condition Variables



Dining Philosophers Example

```
monitor dp
{
    enum {thinking, hungry, eating} state[5];
    condition self[5];
    void pickup(int i)           // following slides
    void putdown(int i)         // following slides
    void test(int i)            // following slides
    void init() {
        for (int i = 0; i < 5; i++)
            state[i] = thinking;
    }
}
```

Dining Philosophers

```
void pickup(int i) {  
    state[i] = hungry;  
    test(i);  
    if (state[i] != eating)  
        self[i].wait();  
}  
  
void putdown(int i) {  
    state[i] = thinking;  
    // test left and right neighbors  
    test((i+4) % 5);  
    test((i+1) % 5);  
}
```

Dining Philosophers

```
void test(int i) {  
    if ( (state[(i + 4) % 5] != eating) &&  
        (state[i] == hungry) &&  
        (state[(i + 1) % 5] != eating)) {  
        state[i] = eating;  
        self[i].signal();  
    }  
}
```

Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
```

- Each external procedure F will be replaced by

```
wait(mutex);
```

```
...
```

```
body of  $F$ ;
```

```
...
```

```
if (next-count > 0)
```

```
    signal(next)
```

```
else
```

```
    signal(mutex);
```

- Mutual exclusion within a monitor is ensured.

Monitor Implementation

- For each condition variable **x**, we have:
 semaphore x-sem; // (initially = 0)
 int x-count = 0;
- The operation **x.wait** can be implemented as:

```
x-count++;  
if (next-count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x-sem);  
x-count--;
```

Monitor Implementation

- The operation **x.signal** can be implemented as:

```
if (x-count > 0) {  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```

Monitor Implementation

■ *Conditional-wait* construct: **x.wait(c);**

- ✦ **c** – integer expression evaluated when the **wait** operation is executed.
- ✦ value of **c** (a *priority number*) stored with the name of the process that is suspended.
- ✦ when **x.signal** is executed, process with smallest associated priority number is resumed next.

■ Check two conditions to establish correctness of system:

- ✦ User processes must always make their calls on the monitor in a correct sequence.
- ✦ Must ensure that an uncooperative process does not ignore the mutual-exclusion gateway provided by the monitor, and try to access the shared resource directly, without using the access protocols.

Solaris 2 Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing.
- Uses *adaptive mutexes* for efficiency when protecting data from short code segments.
- Uses *condition variables* and *readers-writers* locks when longer sections of code need access to data.
- Uses *turnstiles* to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock.

Windows 2000 Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems.
- Uses *spinlocks* on multiprocessor systems.
- Also provides *dispatcher objects* which may act as wither mutexes and semaphores.
- Dispatcher objects may also provide *events*. An event acts much like a condition variable.