

NETWORK DISTRIBUTED SYSTEMS

FAILURES & DISTRIBUTED CONSENSUS

Dr. Christina Thorpe

THE PLAYERS

- Choose from a large set of interchangeable terms:
 - Processes, threads, tasks,...
 - Processors, nodes, servers, clients,...
 - Actors, agents, participants, partners, cohorts...
- The term “node” or “actor” is also popular
 - Short and sweet
 - A logical/virtual entity: may be multiple logical nodes per physical machine.
 - General with regard to role and internal structure
 - Tend to use “actor” if self-interest is an issue

PROPERTIES OF NODES/ ACTORS

Essential properties typically assumed by model:

- Private state
 - Distributed memory: model sharing as messages
- Executes a sequence of state transitions
 - Some transitions are reactions to messages
 - May have internal concurrency, but hide that
- Deterministic vs. nondeterministic
- Unique identity vs. anonymous nodes
- Local clocks with arbitrary drift vs. global time strobe (e.g., GPS satellites)

NODE FAULTS AND FAILURES

- **Fail-stop:** Nodes/actors may fail by stopping.
- **Byzantine:** Nodes/actors may fail without stopping.
 - Arbitrary, erratic, unexpected behaviour
 - May be malicious and disruptive
- **Unfaithful behaviour:**
 - Actors may behave unfaithfully from self-interest.
 - If it is rational, then it is expected.
 - If it is expected, then we can control it.
 - Design in incentives for faithful behaviour, or disincentives for unfaithful behaviour.

NODE RECOVERY

- Fail-stopped nodes may revive/restart
 - Retain identity
 - Lose messages sent to them while failed
 - Arbitrary time to restart...or maybe never
- Restarted node may recover state at time of failure.
 - Lose state in volatile (primary) memory.
 - Restore state in non-volatile (secondary) memory.
 - Writes to non-volatile memory are expensive.
 - Design problem: recover complete states reliably, with minimal write cost.

MESSAGES

- Processes communicate by sending messages.
- Unicast typically assumed
 - Build multicast/broadcast on top
- Use unique process identity (pid) as destination.
- Optional: cryptography
 - (optional) Sender is authenticated.
 - (optional) Message integrity is assured.
 - E.g., using digital signatures or Message Authentication Codes.

DISTRIBUTED SYSTEM MODELS

- **Synchronous model**

- Message delay is bounded and the bound is known.
- E.g., delivery before next tick of a global clock.
- Simplifies distributed algorithms
 - “learn just by watching the clock”
 - absence of a message conveys information.

- **Asynchronous model**

- Message delays are finite, but unbounded/unknown
- More realistic/general than synchronous model.
 - “Beware of any model with stronger assumptions.”
- Strictly harder/weaker than synchronous model.
 - Consensus is not always possible

MESSAGING PROPERTIES

- Other possible properties of the messaging model:
 - Messages may be lost.
 - Messages may be delivered out of order.
 - Messages may be duplicated.
- Do we need to consider these in our distributed system model?
- Or, can we solve them within the asynchronous model, without affecting its foundational properties?
 - E.g., reliable transport protocol such as TCP

THE NETWORK

- Picture a cloud with open unicast and unbounded capacity/bandwidth.
 - Squint and call it the Internet.
- Alternatively, the network could be a graph:
 - Graph models a particular interconnect structure.
 - Examples: star, ring, hypercube, etc.
 - Nodes must forward/route messages.
 - Issues: cut-through, buffer scheduling, etc.
 - Bounded links, blocking send: may deadlock.


STANDARD ASSUMPTIONS

- For this module, we make reasonable assumptions for general Internet systems:
 - Nodes with local state and (mostly) local clocks
 - Asynchronous model: unbounded delay but no loss
 - Fail-stop or Byzantine
 - Node identity with (optional) authentication
 - Allows message integrity
 - No communication-induced deadlock.
 - Can deadlock occur? How to avoid it?
 - Temporary network interruptions are possible.
 - Including partitions

COORDINATION

- If the solution to availability and scalability is to decentralise and replicate functions and data, how do we coordinate the nodes?
 - data consistency
 - update propagation
 - mutual exclusion
 - consistent global states
 - group membership
 - group communication
 - event ordering
 - distributed consensus
 - quorum consensus

CONSENSUS



Generalizes to N nodes/processes.

PROPERTIES FOR CORRECT CONSENSUS

Termination

Every correct process decides some value.

Validity

If all processes propose the same value v , then all correct processes decide v .

Integrity

Every correct process decides at most one value, and if it decides some value v , then v must have been proposed by some process.


Agreement

Every correct process must agree on the same value.

PROPERTIES OF DISTRIBUTED ALGORITHMS

- Agreement is a safety property.
 - Every possible state of the system has this property in all possible executions.
 - I.e., either they have not agreed yet, or they all agreed on the same value.
- Termination is a liveness property.
 - Some state of the system has this property in all possible executions.
 - The property is stable: once some state of an execution has the property, all subsequent states also have it.

VARIANT I: CONSENSUS (C)




$$d_i = v_k$$

P_i selects d_i from $\{v_0, \dots, v_{N-1}\}$.

All P_i select d_i as the same v_k .

If all P_i propose the same v , then $d_i = v$, else d_i is arbitrary.

VARIANT II: COMMAND CONSENSUS (BG)




P_i selects $d_i = v_{\text{leader}}$ proposed by designated **leader** node P_{leader} if the leader is correct, else the selected value is arbitrary.

As used in the *Byzantine generals* problem.

Also called *attacking armies*.

VARIANT III: INTERACTIVE CONSISTENCY (IC)



P_i selects $d_i = [v_0, \dots, v_{N-1}]$ vector reflecting the values proposed by all correct participants.

FISCHER-LYNCH-PATTERSON (1985)

- No consensus can be **guaranteed** in an asynchronous communication system in the presence of any failures.
- Intuition: a “failed” process may just be slow, and can rise from the dead at exactly the wrong time.
- Consensus **may** occur recognisably, rarely or often.
 - e.g., if no inconveniently delayed messages
- FLP implies that no agreement can be guaranteed in an asynchronous system with Byzantine failures either. (More on that later.)

CONSENSUS IN PRACTICE I

- What do these results mean in an asynchronous world?
 - Unfortunately, the Internet is asynchronous, even if we believe that all faults are eventually repaired.
 - Synchronized clocks and predictable execution times don't change this essential fact.
- Even a single faulty process can prevent consensus.
- The FLP impossibility result extends to:
 - Reliable ordered multicast communication in groups
 - Transaction commit for coordinated atomic updates
 - Consistent replication
- These are practical necessities, so what are we to do?

CONSENSUS IN PRACTICE II

- We can use some tricks to apply synchronous algorithms:
 - **Fault masking:** assume that failed processes always recover, and reintegrate them into the group.
 - If you haven't heard from a process, wait longer...
 - A round terminates when every expected message is received.
 - **Failure detectors:** construct a failure detector that can determine if a process has failed.
 - A round terminates when every expected message is received, or the failure detector reports that its sender has failed.
- **But:** protocols may block in pathological scenarios, and they may misbehave if a failure detector is wrong.


FAILURE DETECTORS

- How to detect that a member has failed?
 - pings, timeouts, beacons, heartbeats
 - recovery notifications
 - “I was gone for awhile, but now I’m back.”
- Is the failure detector accurate?
- Is the failure detector live (complete)?
- In an asynchronous system, it is possible for a failure detector to be accurate or live, but not both.
 - FLP tells us that it is impossible for an asynchronous system to agree on anything with accuracy and liveness!


FAILURE DETECTORS IN REAL SYSTEMS

- Use a detector that is accurate but not live.
 - “I’m back....hey, did anyone hear me?”
 - Can’t wait forever...
- Use a detector that is live but not accurate.
 - Assume bounded processing delays and delivery times.
 - Timeout with multiple retries detects failure accurately with high probability. Tune it to observed latencies.
 - If a “failed” site turns out to be alive, then restore it or kill it (fencing, fail-silent).
 - Example: leases and leased locks
- What do we assume about communication failures? How much pinging is enough?
What about network partitions?

A NETWORK PARTITION



TWO GENERALS IN PRACTICE



How do banks solve this problem?

Keith Marzullo

COMMITTING DISTRIBUTED TRANSACTIONS

- Transactions may touch data at more than one site.
- Problem: any site may fail or disconnect while a commit for transaction T is in progress.
 - Atomicity says that T does not “partly commit”, i.e., commit at some site and abort at another.
 - Individual sites cannot unilaterally choose to abort T without the agreement of the other sites.
 - If T holds locks at a site S , then S cannot release them until it knows if T committed or aborted.
 - If T has pending updates to data at a site S , then S cannot expose the data until T commits/aborts.

COMMIT IS A CONSENSUS PROBLEM

- If there is more than one site, then the sites must agree to commit or abort.
- Sites (Resource Managers or RMs) manage their own data, but coordinate commit/abort with other sites.
 - “Log locally, commit globally.”
- We need a protocol for distributed commit.
 - It must be safe, even if FLP tells us it might not terminate.
- Each transaction commit is led by a coordinator (Transaction Manager or TM).